JAMES MORRISON

Data Engineer | Senior Level

Edinburgh, United Kingdom

Phone: +44 7891 234567

Email: james.morrison@gmail.com

LinkedIn: linkedin.com/in/jamesmorrison-dataengineer

PROFESSIONAL SUMMARY

I'm a data engineer with 8 years in the field, though my journey has taken some unconventional turns. Started as a graduate building ETL pipelines, then somehow ended up leading a 45-person data infrastructure team at 26 years old, which taught me a lot about organizational chaos. Now settled into senior-level work where I actually focus on the technical depth I initially wanted. I'm most comfortable with distributed systems, pipeline architecture, and untangling legacy data infrastructure that nobody wants to touch.

WORK EXPERIENCE

Graduate Data Engineer

TechVault Solutions

techvault.io

London, United Kingdom

July 2016 - September 2018

Developed and maintained Python-based ETL pipelines processing 2TB of data daily across multiple cloud environments

Managed direct team of 23 engineers and coordinated departmental roadmap for data infrastructure initiatives

Designed schema optimization strategies that reduced query latency by 40% for analytics queries

Implemented CI/CD pipelines for data workflows using Jenkins and improved deployment frequency from monthly to weekly releases

Senior Infrastructure Architect

CloudNine Analytics Ltd

cloudnine-analytics.co

Glasgow, United Kingdom

October 2018 - March 2021

Architected real-time data streaming solutions using Apache Kafka processing over 500 million events daily

Oversaw infrastructure costs and negotiated with vendors, responsible for £2.1M annual budget allocation

Built data warehouse on Snowflake with dimensional modeling supporting 150+ concurrent analysts

Led migration of legacy Hadoop clusters to cloud-native architecture, reducing operational overhead by 60%

Senior Data Engineer

DataFlow Innovations

dataflowtech.net

Edinburgh, United Kingdom

April 2021 - Present

Designed and implemented scalable data pipelines using Spark and Python handling petabyte-scale datasets across AWS

Optimized Apache Airflow DAG orchestration reducing pipeline execution time by 35% and improving reliability to 99.8%

Developed data quality framework using Great Expectations catching 2000+ anomalies monthly before reaching analytics layer

Mentored 4 junior engineers on distributed systems concepts and best practices for production pipeline development

EDUCATION

BSc Computer Science

University of Edinburgh

Edinburgh, United Kingdom

Graduated: June 2016

SKILLS

Languages and Frameworks: Python, SQL, Scala, PySpark, Apache Spark

Data Platforms and Tools: Apache Kafka, Apache Airflow, Snowflake, AWS (S3, EC2, RDS, Lambda), Google Cloud Platform, Spark SQL

Databases: PostgreSQL, MySQL, DynamoDB, MongoDB

DevOps and Tools: Docker, Kubernetes, Jenkins, Git, Terraform, Linux

Data Concepts: ETL/ELT pipelines, dimensional modeling, data warehousing, real-time streaming, data quality, schema design