JAMES MITCHELL

james.mitchell@gmail.com

+44 7841 923456

linkedin.com/in/james-mitchell-ml-engineer

PROFESSIONAL SUMMARY

I'm a machine learning engineer based in Bristol with seven years of experience building and deploying production ML systems. I've spent the last few years working on computer vision and NLP problems at scale, particularly around recommendation systems and predictive analytics. I'm most interested in taking on leadership responsibilities and mentoring junior engineers while continuing to contribute technically to challenging problems.

WORK EXPERIENCE

Senior Machine Learning Engineer

Nexus Analytics Ltd

nexus-analytics.io

Bristol, United Kingdom

March 2022 - Present

Designed and implemented end-to-end computer vision pipeline for real-time object detection that reduced inference latency by 45% through model optimization and batching strategies

Led migration of legacy Python codebase to production-ready microservices architecture, improving deployment frequency from monthly to weekly releases

Established ML model monitoring and validation framework that caught data drift issues before they impacted business metrics, preventing estimated GBP 200k in losses

Mentored team of three junior engineers on statistical testing, feature engineering best practices, and code review standards

Machine Learning Engineer

DataCore Solutions

datacore.net

Bristol, United Kingdom

June 2019 - February 2022

Developed ensemble learning models for customer churn prediction achieving 88% AUC, directly informing retention strategies that improved customer lifetime value by 12%

Implemented automated feature engineering pipeline using scikit-learn and pandas that reduced manual data preparation time from 4 hours to 20 minutes per project

Built and trained transformer-based NLP models for document classification, processing 500k documents monthly with 94% accuracy

Collaborated with data engineering team to optimize SQL queries and improve data warehouse performance for ML workloads

Machine Learning Engineer

Quantified AI

quantified-ai.co

London, United Kingdom

September 2017 - May 2019

Created predictive maintenance models for industrial equipment using gradient boosting algorithms, helping clients reduce unplanned downtime by average of 30%

Conducted hyperparameter tuning and cross-validation experiments across multiple model architectures to establish performance baselines

Deployed machine learning models to AWS using Flask and Docker, setting up monitoring dashboards and alert systems for model drift

Presented technical findings to non-technical stakeholders quarterly, translating model complexity into business impact and ROI calculations

EDUCATION

MSc Machine Learning

University of Bristol

Bristol, United Kingdom

Graduated September 2017

BSc Computer Science

University of Manchester

Manchester, United Kingdom

Graduated June 2015

TECHNICAL SKILLS

Machine Learning Frameworks: TensorFlow, PyTorch, scikit-learn, XGBoost, LightGBM

Data Processing: Python, pandas, NumPy, SQL, PySpark

Cloud Platforms: AWS (SageMaker, EC2, S3), Google Cloud Platform

Deployment and DevOps: Docker, Kubernetes, Terraform, ArgoCD, CI/CD pipelines

Computer Vision: OpenCV, YOLO, CNN architectures, transfer learning

Natural Language Processing: spaCy, NLTK, transformers, BERT, GPT fine-tuning

Data Visualization: Matplotlib, Seaborn, Plotly, Tableau

Statistical Analysis: hypothesis testing, A/B testing, Bayesian methods, experimentation design

Databases: PostgreSQL, MongoDB, Elasticsearch

Version Control: Git, GitHub, GitLab