

MARCUS FRIEDRICH

Data Engineer

Email: marcus.friedrich@gmail.com

Phone: +49 30 7654 3821

LinkedIn: [linkedin.com/in/marcus-friedrich-berlin](https://www.linkedin.com/in/marcus-friedrich-berlin)

GitHub: github.com/mfriedrich-data

PROFESSIONAL SUMMARY

Results-driven Data Engineer with 4 years of experience designing and maintaining scalable data pipelines and infrastructure. Skilled in building ETL processes, optimizing database performance, and implementing data warehousing solutions. Known for translating complex requirements into robust technical implementations and mentoring junior team members. Passionate about clean code, data quality, and leveraging modern tools to solve real-world business problems.

TECHNICAL SKILLS

Programming Languages: Python, Scala, SQL, Java

Big Data Technologies: Apache Spark, Hadoop, Kafka, Flink

Cloud Platforms: AWS, Google Cloud Platform, Azure

Data Warehousing: Snowflake, Redshift, BigQuery, PostgreSQL

ETL Tools: Airflow, Talend, Apache NiFi

Data Formats: Parquet, Avro, JSON, CSV, Protocol Buffers

Version Control: Git, GitHub, GitLab

Other Tools: Docker, Kubernetes, Tableau, dbt

WORK EXPERIENCE

Principal Data Architect

TechVision Solutions

Munich, Germany

<https://techvision-solutions.io>

January 2023 - Present

Directed the strategic oversight of enterprise data infrastructure serving 300+ data consumers across 15 business units with annual budgets exceeding 5 million euros

Architected and deployed a real-time data platform processing 500 billion events daily using Kafka and Spark, reducing query latency by 78 percent

Managed a global team of 45 data engineers across Berlin, London, and Singapore offices, implementing performance management systems and career development frameworks

Established data governance policies and compliance frameworks ensuring GDPR and ISO 27001 certification across all data operations

Senior Data Engineer

DataFlow Systems GmbH

<https://dataflow-systems.net>

Berlin, Germany

September 2021 - December 2022

Developed and optimized 12 critical ETL pipelines using Apache Airflow, improving data delivery time by 45 percent and reducing operational costs by 320,000 euros annually

Designed a multi-tenant data warehouse architecture in Snowflake supporting 50 concurrent users and 10 terabytes of analytical data

Implemented automated data quality monitoring using Great Expectations, reducing data incidents by 62 percent

Mentored 3 junior data engineers and conducted technical interviews for recruitment of 8 new team members

Junior Data Engineer

CloudData Analytics

<https://clouddata-analytics.co>

Berlin, Germany

June 2020 - August 2021

Created and maintained 8 Python-based ETL scripts processing daily ingestion of 50 gigabytes of customer transaction data from multiple sources into data lakes

Collaborated with analytics team to build Tableau dashboards and reports from processed datasets, supporting executive decision-making

Wrote comprehensive technical documentation and runbooks for data pipeline operations and troubleshooting procedures

Performed ad-hoc SQL analysis supporting product and marketing teams with customer segmentation and behavioral analytics

EDUCATION

Master of Science in Computer Science

Technical University of Munich (TUM)

Munich, Germany

2018 - 2020

Bachelor of Science in Information Technology

Humboldt University of Berlin

Berlin, Germany

2014 - 2018

CERTIFICATIONS

AWS Certified Data Analytics Specialty

Google Cloud Professional Data Engineer

Certified Kubernetes Administrator (CKA)

GITHUB REPOSITORIES

github.com/mfriedrich-data/spark-pipeline-optimizer

Advanced Spark optimization library with custom partitioning strategies and query plan analysis tools

github.com/mfriedrich-data/airflow-data-monitoring

Open-source Airflow plugin for automated data quality checks and anomaly detection in production pipelines

github.com/mfriedrich-data/snowflake-cost-analyzer

Python utility for analyzing and optimizing Snowflake query costs and resource utilization

github.com/mfriedrich-data/kafka-stream-processor

Real-time stream processing framework built on Kafka for event-driven data transformations