

MARKUS VAN DER BERG

markus.vdberg@gmail.com

+31 6 487 392 154

<https://www.linkedin.com/in/markus-van-der-berg>

<https://github.com/mvdberg>

PROFESSIONAL SUMMARY

Data scientist with 8 years of experience building machine learning systems and analytical pipelines for financial services and e-commerce companies. I've led teams through complex data challenges, from initial data quality assessments to production model deployments. My background is in statistical modeling and software engineering, which helps me bridge the gap between research prototypes and reliable systems. I'm based in Amsterdam and enjoy mentoring junior data scientists while staying hands-on with code.

TECHNICAL SKILLS

Machine Learning: Python, scikit-learn, TensorFlow, PyTorch, XGBoost, LightGBM

Data Engineering: SQL, Apache Spark, Pandas, data pipeline development, ETL

Cloud Platforms: AWS (SageMaker, EC2, S3), Google Cloud Platform, Azure

Databases: PostgreSQL, MySQL, MongoDB, Elasticsearch

Tools & Platforms: Git, Docker, Kubernetes, Jupyter, MLflow, Airflow

Statistics: Hypothesis testing, A/B testing, time series analysis, causal inference

Data Visualization: Matplotlib, Seaborn, Plotly, Tableau

WORK EXPERIENCE

Senior Data Scientist

Quantify Analytics

Amsterdam, Netherlands

www.quantifyanalytics.io

March 2021 - Present

Led the development of a customer lifetime value prediction model that improved marketing budget allocation by 23 percent, working directly with the VP of Marketing to define success metrics and iterate on features. Built and maintained an automated machine learning pipeline using Airflow and Python that reduced manual modeling work by 15 hours per week across the team. Established best practices for model validation and monitoring, catching data drift issues before they impacted production predictions. Mentored 2 junior data scientists on feature engineering, statistical testing, and Python development standards.

Data Scientist

FinTech Solutions B.V.

Utrecht, Netherlands

<https://www.fintech-solutions.nl>

June 2018 - February 2021

Developed fraud detection models using gradient boosting that reduced false positives by 31 percent while maintaining detection accuracy, directly improving customer experience. Designed and executed A/B tests for 6 different pricing strategies, providing statistical evidence that informed company-wide pricing decisions generating 2.4 million euros in additional annual revenue. Created a real-time anomaly detection system for transaction monitoring using isolation forests, processing over 100,000 transactions daily. Collaborated with product and

engineering teams to translate complex statistical findings into actionable business recommendations.

Data Analyst

DataDriven Insights

Rotterdam, Netherlands

www.datadriveninsights.co

July 2016 - May 2018

Built ETL pipelines in Python and SQL that consolidated data from 12 different source systems into a unified analytics warehouse, enabling faster reporting and reducing data reconciliation time by 40 percent. Performed statistical analysis on customer behavior datasets to identify high-value segments, leading to targeted campaigns that increased conversion rates from 2.1 percent to 3.7 percent. Developed interactive dashboards in Tableau for executive leadership tracking key performance indicators across 8 business units. Cleaned and prepared datasets for analysis, working with inconsistent and missing data sources while documenting data quality issues.

EDUCATION

Master of Science in Data Science

University of Amsterdam

Amsterdam, Netherlands

Graduated: August 2016

Thesis: Predicting Customer Churn in E-commerce Using Deep Learning Methods

Bachelor of Science in Mathematics

University of Amsterdam

Amsterdam, Netherlands

Graduated: June 2014

GITHUB REPOSITORIES

<https://github.com/mvdberg/churn-prediction-models>

Classification models for predicting customer churn with comparative analysis of logistic regression, random forests, and gradient boosting approaches.

<https://github.com/mvdberg/time-series-forecasting>

Implementation of ARIMA, Prophet, and LSTM models for demand forecasting with real retail data and performance benchmarks.

<https://github.com/mvdberg/feature-engineering-toolkit>

Reusable Python library for automated feature engineering including temporal aggregations, lag features, and categorical encoding methods.

<https://github.com/mvdberg/ml-monitoring-dashboard>

Flask-based monitoring system for tracking model performance metrics, data drift, and prediction distribution changes in production environments.