MARCUS WEBER

Zurich, Switzerland

Email: marcus.weber@gmail.com

Phone: +41 79 245 8932

LinkedIn: https://www.linkedin.com/in/marcus-weber-backend

GitHub: https://github.com/marcusweber

PROFESSIONAL SUMMARY

Backend engineer with 4 years of experience building scalable APIs and microservices for fintech and e-commerce platforms. I've spent the last two years working on high-throughput payment processing systems and recently led the migration of a legacy monolith to a Kubernetes-based architecture. I'm passionate about writing clean code and mentoring junior developers on my team. When I'm not debugging, I contribute to open-source projects and tinker with database optimization problems.

WORK EXPERIENCE

Senior Backend Engineer

FinFlow Systems, Zurich, Switzerland

https://www.finflow.io

January 2023 - Present

- Architected and deployed a new payment reconciliation microservice handling 10 million transactions daily, reducing processing latency by 40% through optimized database queries and Redis caching strategies
- Led the migration of a Python monolith to Go microservices, mentoring three junior developers through the transition and establishing coding standards that reduced code review times by 25%
- Implemented comprehensive monitoring and alerting using Prometheus and Grafana, catching production issues before customer impact in 95% of cases
- Designed and executed database optimization project that reduced PostgreSQL query times by 60%, saving the company approximately 30% on infrastructure costs

Backend Engineer

DataCore Technologies, Zurich, Switzerland

https://www.datacore.net

June 2021 - December 2022

- Developed and maintained 12 REST APIs serving mobile and web clients, managing a combined traffic of 2 million requests per day with 99.8% uptime
- Implemented JWT-based authentication system and role-based access control, improving security audit compliance from 73% to 100%
- Built automated data pipeline using Apache Kafka to process clickstream events, enabling the analytics team to generate real-time reports previously only available after 24-hour delays
- Collaborated with frontend team to design and document GraphQL schema, reducing API endpoints from 47 to 18 and simplifying client integration

Junior Backend Engineer

CloudVault Solutions, Zurich, Switzerland

https://www.cloudvault.co

August 2020 - May 2021

- Developed backend services in Node.js and Express, contributing to the core API that powered the company's document storage platform used by 50,000 active users

- Fixed critical security vulnerability related to file upload validation that could have allowed arbitrary file execution, implemented additional validation layers and wrote comprehensive test coverage

- Participated in on-call rotation and responded to production incidents, gaining hands-on experience with debugging distributed systems and writing post-mortem analyses

- Built internal tools for database management and deployment automation using Python, saving the operations team approximately 5 hours per week in manual tasks

EDUCATION

Bachelor of Science in Computer Science

ETH Zurich (Swiss Federal Institute of Technology Zurich)

Zurich, Switzerland

Graduated: June 2020

TECHNICAL SKILLS

Languages: Go, Python, Node.js, JavaScript, SQL

Databases: PostgreSQL, MySQL, MongoDB, Redis

Message Queues: Apache Kafka, RabbitMQ, AWS SQS

Cloud and DevOps: Docker, Kubernetes, AWS (EC2, RDS, S3, Lambda), GitLab CI/CD

APIs and Tools: REST, GraphQL, gRPC, Git, Linux, Postman

Other: Microservices Architecture, System Design, Database Optimization, Unix/Linux

PUBLIC REPOSITORIES

https://github.com/marcusweber/fast-api-starter

A production-ready Python FastAPI boilerplate with Docker, PostgreSQL integration, and JWT authentication. 87 stars.

https://github.com/marcusweber/kafka-batch-processor

Go library for efficient batch processing of Kafka messages with configurable batching and retry logic. 34 stars.

https://github.com/marcusweber/postgres-query-optimizer

Python tool for analyzing PostgreSQL slow queries and suggesting index optimizations. 121 stars.

https://github.com/marcusweber/distributed-cache-benchmark

Benchmarking suite comparing Redis, Memcached, and DragonflyDB performance under various workloads. 56 stars.