# JAMES MITCHELL

james.mitchell@gmail.com

+44 7891 234567

linkedin.com/in/jamesmitchell-ml

## PROFESSIONAL SUMMARY

Machine Learning Engineer with 6 years of hands-on experience building production systems and deploying models at scale. I've spent the last few years moving between startups and established tech companies, focusing on computer vision and recommendation systems. I'm most interested in solving real-world problems where ML can make a tangible difference, rather than chasing metrics. Outside of work, I enjoy contributing to open-source projects and mentoring junior engineers.

## WORK EXPERIENCE

Senior Machine Learning Engineer

DataFlow Systems

dataflow-systems.io

Cambridge, United Kingdom

March 2022 - Present

Designed and deployed a real-time recommendation engine serving 2 million daily active users, reducing recommendation latency from 400ms to 120ms through model optimization and caching strategies. Led a team of three junior engineers in building computer vision pipelines for product classification, achieving 96.2% accuracy on internal test sets. Established MLOps best practices including automated model evaluation, A/B testing frameworks, and monitoring dashboards that reduced model debugging time by 40%. Collaborated with product and data teams to define success metrics and iterate on features based on user feedback and performance data.

Machine Learning Engineer

Vertex AI Solutions

vertex-ai.co

London, United Kingdom

July 2019 - February 2022

Developed and maintained machine learning models for anomaly detection in financial transactions, processing over 500 million data points daily with 99.2% uptime. Built feature engineering pipelines using PySpark that reduced data processing time from 8 hours to 45 minutes, enabling daily model retraining. Implemented hyperparameter tuning infrastructure using Optuna and Kubernetes, improving model performance by an average of 8% across all deployed models. Conducted code reviews and contributed to internal documentation, helping standardize ML practices across the organization.

Machine Learning Engineer

InnovateTech

innovatetech.net

Cambridge, United Kingdom

September 2017 - June 2019

Created time series forecasting models for demand prediction in supply chain optimization, achieving 92% accuracy on 6-month forecasts for major retail clients. Preprocessed and cleaned complex datasets containing missing values and outliers, working with stakeholders to understand data quality issues and implement validation checks. Deployed three models to production using Docker and deployed them on AWS EC2 instances,

establishing monitoring and alerting systems for model drift detection. Participated in weekly knowledge-sharing sessions and presented findings to non-technical stakeholders, translating technical results into business impact.

## EDUCATION

MSc Machine Learning

University College London

London, United Kingdom

2015 - 2017

BSc Computer Science

University of Cambridge

Cambridge, United Kingdom

2011 - 2015

## TECHNICAL SKILLS

Machine Learning: TensorFlow, PyTorch, scikit-learn, XGBoost, LightGBM, Keras, Optuna, computer vision, NLP, time series forecasting, recommendation systems

Data Processing: Python, PySpark, Pandas, NumPy, SQL, Apache Airflow, Kafka

Deployment and Infrastructure: Docker, Kubernetes, AWS EC2, AWS S3, GitLab CI/CD, Jenkins, model monitoring and evaluation

Data Analysis: Jupyter Notebook, Matplotlib, Seaborn, statistical analysis, A/B testing, experimental design

Development: Git, RESTful APIs, debugging and profiling, code review