

MARTIN ANDERSEN JENSEN

Copenhagen, Denmark

+45 40 28 35 67

martin.andersen.jensen@gmail.com

linkedin.com/in/martin-andersen-jensen

PROFESSIONAL SUMMARY

Data engineer with six years of hands-on experience building scalable data infrastructure and ETL pipelines across multiple industries. Started my career working with batch processing systems and have evolved to designing real-time data architectures that handle hundreds of millions of events daily. I enjoy solving complex data problems and mentoring junior engineers on best practices. Currently focused on optimizing cloud-native data warehousing solutions for enterprise clients.

WORK EXPERIENCE

Senior Data Engineer

DataFlow Systems

www.dataflowsystems.io

Copenhagen, Denmark

March 2022 - Present

Designed and implemented a real-time streaming architecture processing 500 million events daily using Apache Kafka and Spark Streaming, reducing latency from 45 minutes to under 2 minutes for critical business dashboards.

Led the migration of legacy on-premises data warehouse to AWS Redshift, involving schema redesign and ETL replatforming that reduced infrastructure costs by 35 percent while improving query performance by 60 percent.

Mentored a team of three junior data engineers on SQL optimization, data modeling principles, and Python best practices, with two of them promoted to mid-level positions within 18 months.

Established data quality frameworks using Great Expectations and custom monitoring solutions, catching data anomalies before they reached downstream analytics teams and reducing incident response time by 70 percent.

Data Engineer

TechVista Analytics

www.techvista.com

Aarhus, Denmark

July 2020 - February 2022

Built and maintained 40 plus ETL pipelines using Python and Apache Airflow, orchestrating daily data ingestion from 15 different source systems including APIs, databases, and cloud storage.

Optimized existing Hive and Spark jobs that were running inefficiently, improving execution time by an average of 45 percent through query rewriting and partition strategy adjustments.

Collaborated with data science and analytics teams to design fact and dimension tables for a new star schema data mart, enabling self-service reporting capabilities for 50 business users across the organization.

Implemented comprehensive logging and monitoring using ELK stack, providing visibility into pipeline performance and significantly reducing the time spent debugging data issues.

Junior Data Engineer

CloudBase Solutions

www.cloudbase.net

Copenhagen, Denmark

June 2018 - June 2020

Developed Python scripts to extract and transform data from multiple relational databases and third-party APIs, loading results into PostgreSQL and MySQL data warehouses for downstream analytics.

Assisted in building initial data infrastructure on Google Cloud Platform, including BigQuery dataset design and Cloud Dataflow job creation for processing unstructured log data.

Created SQL views and stored procedures to support reporting requirements, working closely with business analysts to understand data requirements and validate output.

Participated in on-call rotation for data pipeline monitoring and troubleshooting, responding to alerts and resolving issues to ensure 99.2 percent uptime of critical data feeds.

EDUCATION

Master of Science in Computer Science

Technical University of Denmark, Kongens Lyngby, Denmark

Graduated June 2018

Thesis: Distributed Processing Frameworks for Large Scale Data Analytics

Bachelor of Science in Mathematics with Computer Science

University of Copenhagen, Copenhagen, Denmark

Graduated June 2016

TECHNICAL SKILLS

Languages: Python, SQL, Scala, Bash

Big Data and Processing: Apache Spark, Apache Kafka, Apache Airflow, Hive, MapReduce, Flink

Cloud Platforms: AWS (EC2, S3, Redshift, RDS, Lambda), Google Cloud Platform (BigQuery, Dataflow, Pub/Sub), Azure (Data Factory, Synapse)

Databases: PostgreSQL, MySQL, Cassandra, MongoDB, Snowflake

Data Warehousing: Dimensional modeling, fact and dimension tables, slowly changing dimensions, Redshift, BigQuery, Snowflake

ETL and Data Integration: Talend, Informatica, custom Python frameworks

Tools and Frameworks: Git, Docker, Kubernetes, Jenkins, Jupyter, Pandas, NumPy

Data Quality and Monitoring: Great Expectations, dbt, Datadog, CloudWatch, ELK Stack

Methodologies: Agile, Git workflows, code review processes, test-driven development